

✓ APS360 Team42 — News Summarization + Image Retrieval → PDF (Full Pipeline)

Purpose. This notebook turns a cleaned news CSV into a polished multi-page PDF with **model-generated summaries** and **relevant images**. It uses the Team42 systems built for APS360 (UofT, Summer 2025). The workflow is reproducible, GPU-friendly, and includes quantitative evaluation.

Configuration used in this notebook (current run)

- **Summarization:** fine-tuned **BART** loaded from `SUMMARIZER_DIR`.
- **Images: retrieval (default)** — not T2I.
 - **Stage-1 recall:** OpenCLIP **ViT-B/32** text→image over a local indexed corpus (512-D).
 - **Stage-2 rerank:** OpenCLIP **ViT-L/14** to sort candidates.
 - We keep **3 images per article** using a soft similarity floor + “top-up” (never fewer than 3).
- **PDF build:** one page per article (source preview, summary, images).
- **Evaluation:** ROUGE-L, BERTScore, CLIP alignment, duplicate rate, and word-length stats (article vs. summary).
- *(Optional but disabled here)* **Text-to-Image (T2I)** via Diffusers (`IMAGE_MODEL_DIR`), We didn't use it due to a lack of time.

Inputs on Drive

- **Data (CSV):** cleaned news dataset
Link: [Data](#)
- **Models:**
 - **Summarizer directory (`SUMMARIZER_DIR`)** — fine-tuned BART (e.g., `model.safetensors`, `tokenizer.json`, `config.json`, `generation_config.json`, etc.).
 - **(Optional) T2I directory (`IMAGE_MODEL_DIR`)** — a Diffusers pipeline for switching to image generation.

◦ *Link:* [*Models](#)

- **Image Corpus:** `clip_images.rar` (or `.zip`) with JPG/PNG/WEBP.
(*Optional*) `clip_image_index.pt` (prebuilt 512-D image features) to skip re-indexing.

How it works (high level)

1. **Mount & Paths** — set `DATA_CSV_PATH`, `SUMMARIZER_DIR`, `OUTPUTS_DIR` (and optional `IMAGE_MODEL_DIR`).
2. **Load summarizer** — lead-fusion extract→abstract strategy; safe decoding; terminal punctuation cleanup.
3. **Retrieve images** — B/32 for first-stage recall (dimension-matched), L/14 for rerank; keep 3 per article.
4. **Render PDF** — ReportLab layout with header, source preview, summary, and a row of images.
5. **Evaluate** — print ROUGE-L, BERTScore, CLIP alignment, duplication rate, and article vs. summary lengths; preview one summary + its 3 images.

Typical results (our final run)

- ROUGE-L \approx **0.43**, BERTScore \approx **0.32**, CLIP align \approx **0.22**.
- **3 images/article**, summaries \approx **65 words**, near-zero repetition.
- (Values vary with dataset/corpus content.)*

Quick toggles

- `N_ROWS` — how many articles to process.
- `FINAL_K = 3`, `ALIGN_FLOOR` — images per article and similarity floor for selection.
- `USE_RETRIEVAL` (default **True**) vs. `USE_T2I` (set **True** only if `IMAGE_MODEL_DIR` is present).
- Rerank backbone: switch `'ViT-L-14'` ↔ `'ViT-L-14@336px'` for a small CLIP-align bump.

Repro & notes

- Determinism: `torch.manual_seed(42)` and version printouts are included.
- Tokenizer name collisions avoided (`sum_tok` / `sum_model` vs. CLIP tokenizers).

- Dimension safety: first-stage **must** be **ViT-B/32** (512-D) to match `img_feats`; L/14 only for rerank/eval.

Team & credits

Team42 – APS360 (Summer 2025), University of Toronto.

Built with: Hugging Face Transformers, OpenCLIP, ReportLab, rouge-score, bert-score (and optional Diffusers for T2I).

```
# We remove old conflicting installs
!pip uninstall -y numpy scipy scikit-learn

# We install versions that play nicely together
!pip install --upgrade numpy==1.26.4 scipy==1.14.1 scikit-learn==1.5.2

# We reinstall transformers after fixing dependencies
!pip install --upgrade transformers accelerate sentencepiece
```

```
➡ Found existing installation: numpy 1.26.4
Uninstalling numpy-1.26.4:
  Successfully uninstalled numpy-1.26.4
Found existing installation: scipy 1.14.1
Uninstalling scipy-1.14.1:
  Successfully uninstalled scipy-1.14.1
Found existing installation: scikit-learn 1.5.2
Uninstalling scikit-learn-1.5.2:
  Successfully uninstalled scikit-learn-1.5.2
Collecting numpy==1.26.4
  Using cached numpy-1.26.4-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (16.8 MB)
Collecting scipy==1.14.1
  Using cached scipy-1.14.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (39.4 MB)
Collecting scikit-learn==1.5.2
  Using cached scikit_learn-1.5.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (11.1 MB)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist (3.5.0)
Using cached numpy-1.26.4-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (16.8 MB)
Using cached scipy-1.14.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (39.4 MB)
Using cached scikit_learn-1.5.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (11.1 MB)
Installing collected packages: numpy, scipy, scikit-learn
ERROR: pip's dependency resolver does not currently take into account all the
opencv-contrib-python 4.12.0.88 requires numpy<2.3.0,>=2; python_version >= "3.9" but you have numpy 1.26.4 which is incompatible
umap-learn 0.5.9.post2 requires scikit-learn>=1.6, but you have scikit-learn 1.5.2 which is incompatible
opencv-python-headless 4.12.0.88 requires numpy<2.3.0,>=2; python_version >= "3.9" but you have numpy 1.26.4 which is incompatible
thinc 8.3.6 requires numpy<3.0.0,>=2.0.0, but you have numpy 1.26.4 which is incompatible
opencv-python 4.12.0.88 requires numpy<2.3.0,>=2; python_version >= "3.9" but you have numpy 1.26.4 which is incompatible
```

```

OpenCV Python 4.12.0.80 requires numpy>=1.17.0; python_version < 3.10, but
Successfully installed numpy-1.26.4 scikit-learn-1.5.2 scipy-1.14.1
Requirement already satisfied: transformers in /usr/local/lib/python3.11/dist-
Collecting transformers
Using cached transformers-4.55.2-py3-none-any.whl.metadata (41 kB)
Requirement already satisfied: accelerate in /usr/local/lib/python3.11/dist-pa
Requirement already satisfied: sentencepiece in /usr/local/lib/python3.11/dist
Collecting sentencepiece
Using cached sentencepiece-0.2.1-cp311-cp311-manylinux_2_27_x86_64.manylinux
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-pack
Requirement already satisfied: huggingface-hub<1.0,>=0.34.0 in /usr/local/lib/
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.11/dist-p
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/di
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-p
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.11/
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-pack
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/pythor
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.11
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.11/dist-pa
Requirement already satisfied: psutil in /usr/local/lib/python3.11/dist-packag
Requirement already satisfied: torch>=2.0.0 in /usr/local/lib/python3.11/dist-
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/c
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/py
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-pack
Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packag
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch>=2.0.0->accelerate)
Using cached nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.w
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch>=2.0.0->accelerate)
Using cached nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch>=2.0.0->accelerate)
Using cached nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.w
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch>=2.0.0->accelerate)
Using cached nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.me
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch>=2.0.0->accelerate)
Using cached nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.n
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch>=2.0.0->accelerate)
Using cached nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.me
Collecting nvidia-curand-cu12==10.3.5.147 (from torch>=2.0.0->accelerate)
Using cached nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch>=2.0.0->accelerate)
Using cached nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl
Collecting nvidia-cusparselt-cu12==12.3.1.170 (from torch>=2.0.0->accelerate)
Using cached nvidia_cusparselt_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.w
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /usr/local/li
Collecting nvidia-nccl-cu12==2.21.5 (from torch>=2.0.0->accelerate)
Using cached nvidia_nccl_cu12-2.21.5-py3-none-manylinux2014_x86_64.whl.metac
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/py
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch>=2.0.0->accelerate)
Using cached nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.wh
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist

```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
#All the models and data are saved in our drive
```

```
DRIVE_BASE      = "/content/drive/MyDrive"
SUMMARIZER_DIR  = f"{DRIVE_BASE}/model"
IMAGE_MODEL_DIR = f"{DRIVE_BASE}/image_model"
DATA_CSV_PATH   = f"{DRIVE_BASE}/data/cleaned_dataset.csv"
OUTPUTS_DIR     = f"{DRIVE_BASE}/outputs"
```

```
# Create outputs dir if needed
```

```
import os
os.makedirs(OUTPUTS_DIR, exist_ok=True)
```

```
⇒ Drive already mounted at /content/drive; to attempt to forcibly remount, call
   drive.mount('/content/drive', force_remount=True)
```

```
#Lead-Fusion Extract→Abstract; collision-proof
```

```
import re, torch
from transformers import AutoTokenizer, AutoModelForSeq2SeqLM
```

```
device = "cuda" if torch.cuda.is_available() else "cpu"
torch.manual_seed(42)
```

```
sum_tok = AutoTokenizer.from_pretrained(SUMMARIZER_DIR, use_fast=True)
sum_model = AutoModelForSeq2SeqLM.from_pretrained(
    SUMMARIZER_DIR,
    torch_dtype=torch.float16 if device == "cuda" else torch.float32,
    low_cpu_mem_usage=True,
).to(device)
sum_model.eval()
```

```
TARGET_WORDS    = 72
MIN_WORDS       = 58
MAX_NEW_CAP     = 210 # hard cap in tokens
MIN_NEW_FLOOR   = 80
```

```
# sentence utilities
```

```
_sent_split = re.compile(r'(?<=[.!?])\s+')
_stop = set("""
```

```
a an the and or but if while of in on at by for with without into from to as that
i you he she it we they them his her their our your its not no yes do did does do
```

```

over under between about after before during around up down out off so than then :
"".split())

def _words(s): return [w.lower() for w in re.findall(r"[A-Za-z0-9%$€£\-]+", s)]

def _score_sents(sents, lead_bias=1.7):
    # unigram-salience + strong lead bias
    freqs = {}
    for s in sents:
        for w in _words(s):
            if w not in _stop:
                freqs[w] = freqs.get(w, 0) + 1
    scores = []
    for i, s in enumerate(sents):
        sc = sum(freqs.get(w,0) for w in _words(s) if w not in _stop)
        sc *= (lead_bias ** (-i/3.0))
        scores.append(sc)
    return scores

def _lead_fusion(text, lead_k=3, extra_k=2):
    sents = [s.strip() for s in _sent_split.split(text) if s.strip()]
    if len(sents) <= lead_k:
        return " ".join(sents)
    lead = sents[:lead_k]
    rest = sents[lead_k:]
    if not rest:
        return " ".join(lead)
    scores = _score_sents(rest)
    top_idx = sorted(range(len(rest)), key=lambda i: (-scores[i], i))[:extra_k]
    fused = lead + [rest[i] for i in top_idx]
    return " ".join(fused)

# chunking (fallback for very long)
def _chunk_by_tokens(text: str, max_tokens: int, overlap: int = 160):
    ids = sum_tok.encode(text, add_special_tokens=False)
    if len(ids) <= max_tokens: return [text]
    chunks, i = [], 0
    while i < len(ids):
        window = ids[i:i+max_tokens]
        chunks.append(sum_tok.decode(window, skip_special_tokens=True))
        if i + max_tokens >= len(ids): break
        i += max_tokens - overlap
    return chunks

# decoding tuned for faithfulness (less paraphrase)

```

```

GEN_KW = dict(
    do_sample=False,
    num_beams=4,
    length_penalty=1.00,
    no_repeat_ngram_size=3,
    repetition_penalty=1.00,
    early_stopping=True,
)

def _gen(texts, max_new=180, min_new=90):
    inputs = sum_tok(texts, truncation=True, max_length=1024, return_tensors="pt")
    with torch.inference_mode():
        out = sum_model.generate(
            **inputs,
            max_new_tokens=max_new,
            min_new_tokens=min_new,
            **GEN_KW
        )
    return sum_tok.batch_decode(out, skip_special_tokens=True)

# length targeting around word counts
def _token_targets_for_words(target_words):
    # crude token estimate: ~1.3-1.6x words for encoder-decoder
    return int(target_words * 1.5)

def _length_targets(texts):
    # fixed band for stability across articles
    max_new = min(MAX_NEW_CAP, _token_targets_for_words(TARGET_WORDS))
    min_new = max(MIN_NEW_FLOOR, _token_targets_for_words(MIN_WORDS))
    return max_new, min_new

# cleaner + terminal punctuation
def _clean_summary(s: str):
    s = " ".join(s.split()).replace(" .", ".").replace(" ,", ",")
    parts = [p.strip() for p in _sent_split.split(s) if p.strip()]
    # de-dup near-identical sentences
    out, seen = [], set()
    for p in parts:
        key = re.sub(r'\W+', '', p.lower())
        if key and key not in seen:
            seen.add(key); out.append(p)
    s = " ".join(out)
    if s and s[0].islower(): s = s[0].upper() + s[1:]
    if s and not s.endswith(('.', '!', '?')): s += '.'
    return s

```


```

# summarize (short sources)
def summarize_batch(texts):
    fused = [_lead_fusion(t, lead_k=3, extra_k=2) for t in texts]
    max_new, min_new = _length_targets(fused)
    return [_clean_summary(s) for s in _gen(fused, max_new=max_new, min_new=min_new)]

#ummarize (long sources)
def summarize_long(text: str, chunk_max_tokens=960, chunk_overlap=160):
    base = _lead_fusion(text, lead_k=3, extra_k=3)
    chunks = _chunk_by_tokens(base, max_tokens=chunk_max_tokens, overlap=chunk_overlap)
    if len(chunks) == 1:
        return summarize_batch([chunks[0]])[0]
    first = _gen(chunks, max_new=120, min_new=70)
    merged = " ".join(first)
    max_new, min_new = _length_targets([merged])
    return _clean_summary(_gen([merged], max_new=max_new, min_new=min_new)[0])

# public API
def summarize_any(texts_or_series):
    out = []
    for t in texts_or_series:
        t = str(t or "").strip()
        if not t: out.append(""); continue
        n_tok = len(sum_tok.encode(t, add_special_tokens=False))
        out.append(summarize_batch([t])[0] if n_tok <= 1200 else summarize_long(t))
    return out

```

 /usr/local/lib/python3.11/dist-packages/transformers/models/bart/configuration_bart.py:100: UserWarning: The dtype of the weights is not the same as the dtype of the embedding. This is not a problem for the training, but it might lead to unexpected behavior when using the model for inference.
 warnings.warn(


```

from transformers import AutoConfig
import torch, os

print("Model path:", sum_model.name_or_path)
print("Tokenizer path:", getattr(sum_tok, "name_or_path", "<unknown>"))
print("Device:", next(sum_model.parameters()).device)
print("Param count (M):", round(sum(p.numel() for p in sum_model.parameters())/1e6))

cfg = AutoConfig.from_pretrained(SUMMARIZER_DIR)
print("Config model_type:", cfg.model_type)
print("Local dir exists:", os.path.isdir(SUMMARIZER_DIR))

```

```

➡ Model path: /content/drive/MyDrive/model
Tokenizer path: /content/drive/MyDrive/model
Device: cuda:0
Param count (M): 305.51
Config model_type: bart
Local dir exists: True
/usr/local/lib/python3.11/dist-packages/transformers/models/bart/configuration_bart.py:100:
warnings.warn(

```

```

# --- Block 5: Retrieval-only (disable T2I explicitly) ---
import torch, os

```

```

# Robust device flag
_device = next(sum_model.parameters()).device
is_cuda = (_device.type == "cuda") and torch.cuda.is_available()

```

```

# We are NOT using text-to-image in this run
USE_T2I = False
t2i_pipe = None
print("Text-to-Image disabled; using OpenCLIP retrieval only.")

```

```

# (Optional) sanity: enforce flags other blocks may read
USE_RETRIEVAL = True

```

```

➡ Text-to-Image disabled; using OpenCLIP retrieval only.

```

```

# --- Retrieval fallback: robust + VRAM-safe index build/load ---
!pip -q install "open_clip_torch==2.26.1" ftfy regex tqdm rarfile

```

```

import os, glob, subprocess, zipfile, rarfile, torch
from pathlib import Path
from PIL import Image, UnidentifiedImageError
import open_clip

# Reduce CUDA fragmentation (optional but helpful)
os.environ.setdefault("PYTORCH_CUDA_ALLOC_CONF", "max_split_size_mb:128")

device = "cuda" if torch.cuda.is_available() else "cpu"

# 1) Locate archive on Drive
DRIVE_BASE = "/content/drive/MyDrive"
candidates = glob.glob(f"{DRIVE_BASE}/**/*.clip_images.rar", recursive=True) + \
    glob.glob(f"{DRIVE_BASE}/**/*.clip_images.zip", recursive=True)
if not candidates:
    raise FileNotFoundError("Could not find clip_images.rar (or .zip) anywhere un
ARCHIVE_PATH = candidates[0]
print("Found archive:", ARCHIVE_PATH)

# 2) Extract to /content/clip_images
EXTRACT_DIR = Path("/content/clip_images")
if not EXTRACT_DIR.exists() or not any(EXTRACT_DIR.rglob("*")):
    EXTRACT_DIR.mkdir(parents=True, exist_ok=True)
    if ARCHIVE_PATH.lower().endswith(".rar"):
        try:
            with rarfile.RarFile(ARCHIVE_PATH) as rf:
                rf.extractall(EXTRACT_DIR.as_posix())
        except rarfile.RarCannotExec:
            subprocess.run(["apt-get", "update", "-y"], check=True)
            subprocess.run(["apt-get", "install", "-y", "unrar"], check=True)
            subprocess.run(["unrar", "x", "-o+", ARCHIVE_PATH, EXTRACT_DIR.as_posix()
    else:
        with zipfile.ZipFile(ARCHIVE_PATH) as zf:
            zf.extractall(EXTRACT_DIR.as_posix())

# 3) Collect valid images (skip corrupt)
patterns = ["**/*.jpg", "**/*.jpeg", "**/*.png", "**/*.webp", "**/*.JPG", "**/*.JPEG",
raw_paths = []
for pat in patterns:
    raw_paths.extend(glob.glob(str(EXTRACT_DIR / pat), recursive=True))

image_paths = []
for p in raw_paths:
    try:
        with Image.open(p) as im:

```

```

        im.verify()
        image_paths.append(p)
    except (UnidentifiedImageError, OSError):
        pass

if not image_paths:
    raise FileNotFoundError("No valid jpg/png/webp found in /content/clip_images")
print(f"Images found (valid): {len(image_paths)}")

# 4) Build or load index (CPU tensors for portability/safety)
INDEX_OUT = "/content/clip_image_index.pt"
img_feats = None

if os.path.exists(INDEX_OUT):
    blob = torch.load(INDEX_OUT, map_location="cpu") # load to CPU first
    if isinstance(blob, dict) and "feats" in blob and "paths" in blob and len(blob["feats"]) == len(blob["paths"]):
        img_feats = blob["feats"] # CPU
        image_paths = blob["paths"]
        print("Loaded existing index:", INDEX_OUT, "| Feats shape:", tuple(img_feats.shape))
    else:
        print("Index mismatch → rebuilding ...")
        os.remove(INDEX_OUT)

if img_feats is None:
    # Use quick-gelu variant to match 'openai' tag (removes QuickGELU warning)
    model_name = "ViT-B-32-quickgelu"
    clip_model, _, preprocess = open_clip.create_model_and_transforms(model_name,
        clip_model.eval().to(device)

    # Batch size by VRAM
    if device == "cuda":
        total_mem = torch.cuda.get_device_properties(0).total_memory
        batch = 64 if total_mem >= 20_000_000_000 else 32
    else:
        batch = 16

    feats = []
    with torch.inference_mode():
        for i in range(0, len(image_paths), batch):
            batch_paths = image_paths[i:i+batch]
            ims = []
            for p in batch_paths:
                try:
                    ims.append(preprocess(Image.open(p).convert("RGB")))
                except Exception:

```

```

        # skip late failures
        continue
    if not ims:
        continue
    ims = torch.stack(ims).to(device, non_blocking=True)

    # fp16 on CUDA to cut VRAM
    if device == "cuda":
        with torch.cuda.amp.autocast(dtype=torch.float16):
            f = clip_model.encode_image(ims)
    else:
        f = clip_model.encode_image(ims)

    feats.append(f.detach().float().cpu()) # move to CPU immediately
    del ims, f
    if device == "cuda":
        torch.cuda.empty_cache()

img_feats = torch.cat(feats, dim=0) # CPU tensor
img_feats = img_feats / img_feats.norm(dim=-1, keepdim=True)
torch.save({"feats": img_feats, "paths": image_paths}, INDEX_OUT)
print("Built index:", INDEX_OUT, "| Feats shape:", tuple(img_feats.shape))

# Keep a GPU copy for fast retrieval if available
img_feats = img_feats.to(device, non_blocking=True)

USE_T2I = False
USE_RETRIEVAL = True

# 5) Retrieval function (text -> topk image paths)
def retrieve_images(prompts, topk=2):
    # Ensure a matching text encoder is present
    if 'clip_model' not in globals():
        clip_model_local, _, _ = open_clip.create_model_and_transforms("ViT-B-32-quickgelu")
        clip_model_local.eval().to(device)
    else:
        clip_model_local = clip_model

    tok = open_clip.get_tokenizer("ViT-B-32-quickgelu")
    out = []
    with torch.inference_mode():
        tokens = tok(prompts)
        if isinstance(tokens, dict):
            tokens = {k: v.to(device) for k, v in tokens.items()}
        else:

```

```
print("Retrieval index ready. USE_RETRIEVAL =", USE_RETRIEVAL, "| USE_T2I =", USE_T2I)
```

```
print("USE_T2I =", USE_T2I, "| USE_RETRIEVAL =", USE_RETRIEVAL, "| images:", (len
```

```
# dataset → summaries → images (3 per summary; B/32 retrieve + L/14 rerank)
import pandas as pd, re, numpy as np, os, torch, open_clip
from tqdm import tqdm
from PIL import Image

# data
df = pd.read_csv(DATA_CSV_PATH)
name_map = {c.lower(): c for c in df.columns}
TEXT_COL = (name_map.get("content") or name_map.get("text") or name_map.get("article")
            or name_map.get("body") or name_map.get("document") or name_map.get("story"))
if not TEXT_COL:
    raise ValueError(f"No obvious text column found. Columns: {list(df.columns)}")
```

```

N_ROWS = 30
rows = df[df[TEXT_COL].astype(str).str.strip().ne("")]<div data-bbox="55 188 966 226" data-label="Text">
<div data-bbox="55 265 765 283" data-label="Text">
<div data-bbox="55 303 831 398" data-label="Text">
<div data-bbox="55 437 966 455" data-label="Text">
<div data-bbox="55 475 260 532" data-label="Text">
<div data-bbox="55 552 966 647" data-label="Text">
<div data-bbox="55 667 966 913" data-label="Text">

```

```

        idx = row.topk(topk).indices.tolist()
        out.append([image_paths[i] for i in idx])
    return out  # list[list[path]]

all_images = []

if USE_T2I:
    def make_image_prompts(summary: str):
        return [f"Editorial photo illustrating: {summary}. Global news, realistic
    for s in tqdm(all_summaries, desc="Generating images"):
        ims = generate_images(make_image_prompts(s), num_images=FINAL_K)
        all_images.append(ims)  # list[PIL.Image]

elif USE_RETRIEVAL:
    # 1) First-stage retrieval (B/32)
    queries = [img_query(s) for s in all_summaries]
    first_stage = retrieve_images_b32(queries)  # list[list[path]]

    # 2) Rerank with ViT-L/14 (stronger), apply soft floor, always top-up to FINAL_K
    clipL, _, preprocessL = open_clip.create_model_and_transforms('ViT-L-14', pre
    clipL = clipL.to(device).eval()
    tokL = open_clip.get_tokenizer('ViT-L-14')

    def enc_text_L(txts):
        with torch.inference_mode():
            tks = tokL(txts)
            tks = {k:v.to(device) for k,v in tks.items()} if isinstance(tks, dict)
            t = clipL.encode_text(tks)
            return (t / t.norm(dim=-1, keepdim=True)).float()

    def enc_imgs_L(paths):
        feats, batch = [], 16
        with torch.inference_mode():
            for i in range(0, len(paths), batch):
                ims = []
                for p in paths[i:i+batch]:
                    try: ims.append(preprocessL(Image.open(p).convert("RGB")))
                    except: ims.append(preprocessL(Image.new("RGB", (224,224), (0,0,0))))
                x = torch.stack(ims).to(device)
                feats.append(clipL.encode_image(x))
        v = torch.cat(feats, dim=0)
        return (v / v.norm(dim=-1, keepdim=True)).float()

    text_feats = enc_text_L(all_summaries)

```

```

flat_paths = [p for paths in first_stage for p in paths]
uniq_paths = list(dict.fromkeys(flat_paths))
path_to_idx = {p:i for i,p in enumerate(uniq_paths)}
uniq_feats = enc_imgs_L(uniq_paths)

for i, paths in enumerate(first_stage):
    paths = [p for p in paths if isinstance(p, str) and os.path.exists(p)]
    if not paths:
        all_images.append([])
        continue
    idxs = [path_to_idx[p] for p in paths]
    v = uniq_feats[idxs]                # k x d
    t = text_feats[i:i+1]              # 1 x d
    sims = (t @ v.T).squeeze(0).cpu().numpy() # k
    order = sims.argsort()[::-1]


    picked = [paths[j] for j in order if sims[j] >= ALIGN_FLOOR][:FINAL_K]
    if len(picked) < FINAL_K:           # top-up to 3 regardless
        used = set(picked)
        for j in order:
            if paths[j] not in used:
                picked.append(paths[j]); used.add(paths[j])
            if len(picked) == FINAL_K: break
    if not picked and paths:           # extreme safety
        picked = [paths[order[0]]]

    all_images.append(picked)

else:
    all_images = [[] for _ in all_summaries]

assert len(all_summaries) == len(rows) == len(all_images)
print("Block 5 done:", len(all_summaries), "items; images/item =", np.mean([len(x)
print("sample paths for item 0:", (all_images[0] if all_images else []))


```

 Summarizing: 0%|██████████| 0/4 [00:00<?, ?it/s]Token indices sequence length
Summarizing: 100%|██████████| 4/4 [00:24<00:00, 6.08s/it]
100%|██████████| 933M/933M [00:07<00:00, 117MiB,
Block 5 done: 30 items; images/item = 3.0
sample paths for item 0: ['/content/clip_images/clip_images/wholesaler.jpg',


```
import os, numpy as np
cnts = [len(x) for x in all_images]
exist_ok = [[p for p in x if isinstance(p, str) and os.path.exists(p)] for x in all_images]
print("images/item (mean):", np.mean(cnts), "min:", min(cnts), "max:", max(cnts))
print("missing files per item (first 5):", [len(x)-len(y) for x,y in zip(all_images, exist_ok)])
print("sample paths for item 0:", all_images[0] if all_images else [])
```

```
→ images/item (mean): 3.0 min: 3 max: 3
missing files per item (first 5): [0, 0, 0, 0, 0]
sample paths for item 0: ['/content/clip_images/clip_images/wholesaler.jpg',
```

```
pip install -U reportlab
```

```
→ Collecting reportlab
  Downloading reportlab-4.4.3-py3-none-any.whl.metadata (1.7 kB)
  Requirement already satisfied: pillow>=9.0.0 in /usr/local/lib/python3.11/dist-packages (from reportlab)
  Requirement already satisfied: charset-normalizer in /usr/local/lib/python3.11/dist-packages (from reportlab)
  Downloading reportlab-4.4.3-py3-none-any.whl (2.0 MB)
  ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 2.0/2.0 MB 66.8 MB/s eta 0:00:00
Installing collected packages: reportlab
Successfully installed reportlab-4.4.3
```

```
#Polished PDF builder
from reportlab.lib.pagesizes import A4
from reportlab.pdfgen import canvas
from reportlab.lib.units import cm
from reportlab.lib.styles import getSampleStyleSheet
from reportlab.platypus import Paragraph
from reportlab.lib.enums import TA_LEFT
from reportlab.lib import utils
from textwrap import shorten
from PIL import Image, UnidentifiedImageError
import os
```

```
PAGE_W, PAGE_H = A4
MARGIN = 1.2*cm
CONTENT_W = PAGE_W - 2*MARGIN
```

```
def _wrap_paragraph(text, font_size=11, leading=14):
    styles = getSampleStyleSheet()
    st = styles["Normal"]
```

```

    st.fontName = "Helvetica"
    st.fontSize = font_size
    st.leading = leading
    st.alignment = TA_LEFT
    return Paragraph(text.replace("\n", " ").strip(), st)

def _draw_paragraph(c, p: Paragraph, x, y, max_w):
    w, h = p.wrap(max_w, PAGE_H)
    p.drawOn(c, x, y - h)
    return h

def _pil_to_tmp(img: Image.Image, max_w_px=1600):
    ratio = min(1.0, max_w_px / max(img.size))
    if ratio < 1.0:
        img = img.resize((int(img.width*ratio), int(img.height*ratio)), Image.LANCZOS)
    path = f"/content/_tmp_{id(img)}.jpg"
    img.save(path, "JPEG", quality=88, optimize=True)
    return path

def _draw_image_row_pils(c, imgs, x, y, max_w, height=7*cm, gap=0.5*cm):
    # imgs: list[PIL.Image]
    imgs = imgs[:3]
    if not imgs: return 0
    cols = len(imgs)
    cell_w = (max_w - (cols-1)*gap)/cols
    max_h = 0
    for ci, im in enumerate(imgs):
        path = _pil_to_tmp(im)
        iw, ih = utils.ImageReader(path).getSize()
        scale = min(cell_w/iw, height/ih)
        w, h = iw*scale, ih*scale
        xi = x + ci*(cell_w+gap) + (cell_w - w)/2
        yi = y - h
        c.drawImage(path, xi, yi, width=w, height=h)
        max_h = max(max_h, h)
    return max_h

def _draw_image_row_paths(c, paths, x, y, max_w, height=7*cm, gap=0.5*cm):
    # paths: list[str]
    paths = [p for p in paths if isinstance(p, str) and os.path.exists(p)]
    paths = paths[:3]
    if not paths: return 0
    cols = len(paths)
    cell_w = (max_w - (cols-1)*gap)/cols
    max_h = 0

```

```

for ci, p in enumerate(paths):
    try:
        iw, ih = utils.ImageReader(p).getSize()
    except Exception:
        # Try reopen/convert via PIL as fallback
        try:
            im = Image.open(p).convert("RGB")
            p = _pil_to_tmp(im)
            iw, ih = utils.ImageReader(p).getSize()
        except (UnidentifiedImageError, FileNotFoundError):
            continue
    scale = min(cell_w/iw, height/ih)
    w, h = iw*scale, ih*scale
    xi = x + ci*(cell_w+gap) + (cell_w - w)/2
    yi = y - h
    c.drawImage(p, xi, yi, width=w, height=h)
    max_h = max(max_h, h)
return max_h

def _footer(c, idx):
    c.setFont("Helvetica", 9)
    c.drawRightString(PAGE_W - MARGIN, 0.8*cm, f"Page {idx}")

def build_pdf(sources, summaries, images_per_item, out_path):
    c = canvas.Canvas(out_path, pagesize=A4)
    for idx, (src, summ, imgs) in enumerate(zip(sources, summaries, images_per_item)):
        c.setTitle("Automated News Summary & Images")
        c.setAuthor("Your Pipeline")

        # Header
        c.setFont("Helvetica-Bold", 15)
        c.drawString(MARGIN, PAGE_H - MARGIN, f"News Summary {idx}")

        # Source (truncated preview)
        preview = shorten(str(src), width=600, placeholder="...")
        p1 = _wrap_paragraph(f"<b>Source (preview):</b> {preview}", font_size=10,
            top_y = PAGE_H - MARGIN - 1.0*cm)
        h1 = _draw_paragraph(c, p1, MARGIN, top_y, CONTENT_W)

        # Summary
        p2 = _wrap_paragraph(f"<b>Model Summary:</b> {summ}", font_size=12, leading=1.2)
        h2 = _draw_paragraph(c, p2, MARGIN, top_y - h1 - 0.6*cm, CONTENT_W)

        # Images
        y_start = top_y - h1 - h2 - 1.0*cm

```

```

        if isinstance(imgs, list) and imgs:
            if isinstance(imgs[0], Image.Image):
                _draw_image_row_pils(c, imgs, MARGIN, y_start, CONTENT_W, height=
            elif isinstance(imgs[0], str):
                _draw_image_row_paths(c, imgs, MARGIN, y_start, CONTENT_W, height=

        _footer(c, idx)
        c.showPage()
    c.save()

```

```

# Build the PDF
SRC_PREV_COL = TEXT_COL
pdf_path = "/content/News_Summary_Deck.pdf"
build_pdf(rows[SRC_PREV_COL].tolist(), all_summaries, all_images, pdf_path)

```

```

# Copy to Drive
import shutil, os
os.makedirs(OUTPUTS_DIR, exist_ok=True)
final_path = f"{OUTPUTS_DIR}/News_Summary_Deck.pdf"
shutil.copy(pdf_path, final_path)
print("Saved:", final_path)

```

 Saved: /content/drive/MyDrive/outputs/News_Summary_Deck.pdf

```

# Summary & Image Evaluation (ViT-L-14 matched to rerank) + Article length + Samp

```

```

!pip -q install rouge-score bert-score open_clip_torch ftfy regex tqdm imagehash |

```

```

import torch, numpy as np, pandas as pd
from rouge_score import rouge_scorer
import bert_score
import open_clip
from PIL import Image
import imagehash, os
import matplotlib.pyplot as plt

```

```

device = "cuda" if torch.cuda.is_available() else "cpu"

```

```

# Prepare inputs

```

```

sources = rows[TEXT_COL].astype(str).tolist()

```

```

summaries = list(map(str, all_summaries))

# Normalize images to paths
normalized_images = []
for item in all_images:
    if len(item) == 0:
        normalized_images.append([])
    elif isinstance(item[0], Image.Image): # Generated images
        tmp_paths = []
        for im in item:
            p = f"/content/_eval_{id(im)}.jpg"
            im.save(p, "JPEG", quality=90)
            tmp_paths.append(p)
        normalized_images.append(tmp_paths)
    else: # Retrieval (paths)
        normalized_images.append([str(p) for p in item])

print(f"Evaluating {len(sources)} summaries with images.")

```

0) Article vs Summary Length

```

src_len_words = [len(s.split()) for s in sources]
sum_len_words = [len(s.split()) for s in summaries]

print("\nLength (words)")
print(f"articles    : {np.mean(src_len_words):.1f} ± {np.std(src_len_words):.1f} (")
print(f"summaries    : {np.mean(sum_len_words):.1f} ± {np.std(sum_len_words):.1f} (")

```

1) Summary Quality

```

def first_k_sentences(text, k=4):
    s = text.replace("\n", " ")
    parts = [p.strip() for p in s.split('.') if p.strip()]
    return '. '.join(parts[:k]) + ('.' if parts[:k] else '')

refs = [first_k_sentences(s) for s in sources]
scorer = rouge_scorer.RougeScorer(['rougeL'], use_stemmer=True)
rougeL = [scorer.score(r, h)['rougeL'].fmeasure for r, h in zip(refs, summaries)]

P, R, F1 = bert_score.score(cands=summaries, refs=refs, lang='en', rescale_with_b
berts = F1.tolist()

```

```

def repetition_ratio(s: str, ngram=3):

```

```

    toks = s.lower().split()
    if len(toks) < ngram: return 0.0
    ngrams = [' '.join(toks[i:i+ngram]) for i in range(len(toks)-ngram+1)]
    return 1 - (len(set(ngrams)) / len(ngrams))

rep3 = [repetition_ratio(s, 3) for s in summaries]

summary_df = pd.DataFrame({
    "rougeL": rougeL,
    "bertscore": berts,
    "len_words": sum_len_words,
    "rep_trigram": rep3
})

print("\nSummary metrics (mean ± sd)")
for col in ["rougeL", "bertscore", "len_words", "rep_trigram"]:
    print(f"{col:12s}: {summary_df[col].mean():.3f} ± {summary_df[col].std():.3f}")

# 2) Image Alignment & Diversity (ViT-L-14)

clip_model, _, preprocess = open_clip.create_model_and_transforms('ViT-L-14', pre
clip_model = clip_model.to(device)
tok = open_clip.get_tokenizer('ViT-L-14')

def clip_text_feats(texts):
    with torch.inference_mode():
        tokens = tok(texts)
        tokens = {k:v.to(device) for k,v in tokens.items()} if isinstance(tokens,
        t = clip_model.encode_text(tokens)
        return (t / t.norm(dim=-1, keepdim=True)).float()

def clip_img_feats(paths):
    batch, feats = 16, []
    with torch.inference_mode():
        for i in range(0, len(paths), batch):
            ims = []
            for p in paths[i:i+batch]:
                try:
                    im = Image.open(p).convert("RGB")
                    ims.append(preprocess(im))
                except Exception:
                    ims.append(preprocess(Image.new("RGB", (224, 224), (0, 0, 0))))
            ims = torch.stack(ims).to(device)
            feats.append(clip_model.encode_image(ims))

```

```

    f = torch.cat(feats, dim=0)
    return (f / f.norm(dim=-1, keepdim=True)).float()

text_feats = clip_text_feats(summaries)

align_scores = []
dup_rates = []
all_img_paths = [p for paths in normalized_images for p in paths]

if len(all_img_paths) == 0:
    print("\nNo images to evaluate.")
else:
    # Precompute all image embeddings
    img_offsets = []
    cur = 0
    for paths in normalized_images:
        img_offsets.append((cur, cur+len(paths)))
        cur += len(paths)
    all_img_feats = clip_img_feats(all_img_paths)

    for i, paths in enumerate(normalized_images):
        start, end = img_offsets[i]
        if start == end:
            align_scores.append(np.nan)
            dup_rates.append(np.nan)
            continue

        t = text_feats[i:i+1]
        v = all_img_feats[start:end]
        sim = (t @ v.T).squeeze(0).cpu().numpy()
        align_scores.append(float(sim.max()))

    # Duplicate rate
    hashes = [imagehash.phash(Image.open(p).convert("RGB")) for p in paths]
    pairs = 0
    dups = 0
    for a in range(len(hashes)):
        for b in range(a+1, len(hashes)):
            pairs += 1
            if abs(hashes[a] - hashes[b]) <= 4:
                dups += 1
    dup_rates.append(dups / pairs if pairs else 0.0)

image_df = pd.DataFrame({
    "clip_align": align_scores,

```

```

        "dup_rate": dup_rates,
        "images_per_item": [len(p) for p in normalized_images]
    })

    print("\nImage metrics (mean ± sd)")
    for col in ["clip_align", "dup_rate", "images_per_item"]:
        print(f"{col:14s}: {image_df[col].mean():.3f} ± {image_df[col].std():.3f}")

# 3) Show one example summary + its 3 images

if len(all_img_paths) > 0 and len(align_scores) == len(normalized_images):
    try:
        sample_idx = int(np.nanargmax(align_scores))
    except ValueError:
        sample_idx = next((i for i, lst in enumerate(normalized_images) if len(lst) > 0))
    else:
        sample_idx = next((i for i, lst in enumerate(normalized_images) if len(lst) > 0))

print(f"\n=== Sample item #{sample_idx} ===")
print(f"Article length (words): {src_len_words[sample_idx]}")
print(f"Summary length (words): {sum_len_words[sample_idx]}")
print("\nSUMMARY:\n", summaries[sample_idx])

paths = normalized_images[sample_idx][:3]
paths = [p for p in paths if isinstance(p, str) and os.path.exists(p)]
if paths:
    n = len(paths)
    plt.figure(figsize=(4*n, 3))
    for i, p in enumerate(paths, 1):
        try:
            im = Image.open(p).convert("RGB")
        except Exception:
            im = Image.new("RGB", (224,224), (0,0,0))
        ax = plt.subplot(1, n, i)
        ax.imshow(im)
        ax.set_title(f"Image {i}")
        ax.axis("off")
    plt.tight_layout()
    plt.show()
else:
    print("No images found on disk for this item.")

```



Preparing metadata (setup.py) ... done


```
61.1/61.1 kB 4.2 MB/s eta 0:00:00
296.7/296.7 kB 28.1 MB/s eta 0:00:00
```

Building wheel for rouge-score (setup.py) ... done
Evaluating 30 summaries with images.

Length (words)

articles : 692.8 ± 585.2 (median 556.0)

summaries : 65.6 ± 5.5 (median 66.0)

/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: Use
The secret `HF_TOKEN` does not exist in your Colab secrets.

To authenticate with the Hugging Face Hub, create a token in your settings tab
You will be able to reuse this secret in all of your notebooks.

Please note that authentication is recommended but still optional to access private repositories.
warnings.warn(

tokenizer_config.json: 100% 25.0/25.0 [00:00<00:00, 2.89kB/s]

config.json: 100% 482/482 [00:00<00:00, 62.6kB/s]

vocab.json: 100% 899k/899k [00:00<00:00, 35.1MB/s]

merges.txt: 100% 456k/456k [00:00<00:00, 36.4MB/s]

tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 38.3MB/s]

model.safetensors: 100% 1.42G/1.42G [00:06<00:00, 325MB/s]

Some weights of RobertaModel were not initialized from the model checkpoint at
You should probably TRAIN this model on a down-stream task to be able to use it effectively.

Summary metrics (mean ± sd)

rougeL : 0.429 ± 0.148

bertscore : 0.319 ± 0.164

len_words : 65.567 ± 5.600

rep_trigram : 0.000 ± 0.000

Image metrics (mean ± sd)

clip_align : 0.231 ± 0.025

dup_rate : 0.000 ± 0.000

images_per_item: 3.000 ± 0.000

=== Sample item #28 ===

Article length (words): 337

Summary length (words): 70

SUMMARY:

Rupee closes at 75.59 against the dollar, down 70 paise from its previous close.

Image 1



Image 2



Image 3





```
from google.colab import files
import pandas as pd

# Upload the CSV
uploaded = files.upload()
csv_filename = list(uploaded.keys())[0] # Get uploaded file name

# Read the CSV into a DataFrame
df_new = pd.read_csv(csv_filename)

print(f"Loaded {len(df_new)} rows from {csv_filename}")
print("Columns:", df_new.columns.tolist())
```



Choose Files no files selected

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving new_test_data.csv to new_test_data.csv
Loaded 5 rows from new_test_data.csv
Columns: ['url', 'title', 'text', 'summary']

```
# --- Reset and prepare summaries+images for new_test_data.csv ---
import re, os, torch, open_clip
import pandas as pd
from PIL import Image
from tqdm import tqdm
import numpy as np

assert 'sum_model' in globals() and 'sum_tok' in globals(), "Load your fine-tuned
assert 'img_feats' in globals() and 'image_paths' in globals(), "Build the retrie

device = "cuda" if torch.cuda.is_available() else "cpu"

# Load new CSV and pick the text column
DATA_CSV_PATH = "new_test_data.csv"
df = pd.read_csv(DATA_CSV_PATH)
name_map = {c.lower(): c for c in df.columns}
TEXT_COL = (name_map.get("text") or name_map.get("content") or name_map.get("arti
```

```

        or name_map.get("body") or name_map.get("document") or name_map.get("
assert TEXT_COL, f"No obvious text column found. Columns: {list(df.columns)}"

rows = df[df[TEXT_COL].astype(str).str.strip().ne("")]
print("Rows to process:", len(rows))

# Summarize (uses your fine-tuned model)
def summarize_fn(batch_texts, max_new_tokens=180, min_new_tokens=40):
    tok = sum_tok
    model = sum_model
    outs = []
    for txt in batch_texts:
        inputs = tok(txt, return_tensors="pt", truncation=True, max_length=1024).
        gen = model.generate(
            **inputs,
            do_sample=False,
            num_beams=4,
            length_penalty=1.0,
            max_new_tokens=max_new_tokens,
            min_new_tokens=min_new_tokens,
            no_repeat_ngram_size=3,
            eos_token_id=tok.eos_token_id,
        )
        outs.append(tok.decode(gen[0], skip_special_tokens=True).strip())
    return outs

BATCH = 8
all_summaries = []
for i in tqdm(range(0, len(rows), BATCH), desc="Summarizing (new CSV)"):
    batch = rows.loc[i:i+BATCH-1, TEXT_COL].astype(str).tolist()
    all_summaries.extend(summarize_fn(batch))

assert len(all_summaries) == len(rows), "Summaries length mismatch."
print("Summaries ready:", len(all_summaries))

# Retrieval: B/32 quickgelu recall → L/14 rerank (matches index and eval)
img_feats = img_feats.to(device).float()

def img_query(summary: str):
    keys = re.findall(r'\b([A-Z][a-zA-Z]{2,}|[A-Z]{2,5}|\d{4}|USD|EUR|%|[A-Z][a-z
    key_str = " ".join(dict.fromkeys(keys))
    return (summary + " " + key_str).strip()

RETR_TOPK    = 12
FINAL_K      = 3

```

```
ALIGN_FLOOR = 0.22
```

```
def retrieve_images_b32(prompts, topk=RETR_TOPK):
    model_name = 'ViT-B-32-quickgelu'
    enc_model, _, _ = open_clip.create_model_and_transforms(model_name, pretrained='ViT-B-32-quickgelu')
    enc_model = enc_model.to(device).eval()
    enc_tok = open_clip.get_tokenizer(model_name)
    out = []
    with torch.inference_mode():
        tokens = enc_tok(prompts)
        tokens = {k:v.to(device) for k,v in tokens.items()} if isinstance(tokens, dict) else tokens
        if device == "cuda":
            with torch.amp.autocast('cuda', dtype=torch.float16):
                tfeat = enc_model.encode_text(tokens)
        else:
            tfeat = enc_model.encode_text(tokens)
        tfeat = (tfeat / tfeat.norm(dim=-1, keepdim=True)).float()
        sims = tfeat @ img_feats.T
        for row in sims:
            idx = row.topk(topk).indices.tolist()
            out.append([image_paths[i] for i in idx])
    return out
```

```
# First-stage recall
queries = [img_query(s) for s in all_summaries]
first_stage = retrieve_images_b32(queries)
```

```
# Rerank with ViT-L/14
clipL, _, preprocessL = open_clip.create_model_and_transforms('ViT-L-14', pretrained='ViT-L-14')
clipL = clipL.to(device).eval()
tokL = open_clip.get_tokenizer('ViT-L-14')
```

```
def enc_text_L(txts):
    with torch.inference_mode():
        tks = tokL(txts)
        tks = {k:v.to(device) for k,v in tks.items()} if isinstance(tks, dict) else tks
        if device == "cuda":
            with torch.amp.autocast('cuda', dtype=torch.float16):
                t = clipL.encode_text(tks)
        else:
            t = clipL.encode_text(tks)
    return (t / t.norm(dim=-1, keepdim=True)).float()
```

```
def enc_imgs_L(paths):
    feats, batch = [], 16
```

```

with torch.inference_mode():
    for i in range(0, len(paths), batch):
        ims = []
        for p in paths[i:i+batch]:
            try: ims.append(preprocessL(Image.open(p).convert("RGB")))
            except: ims.append(preprocessL(Image.new("RGB", (224,224), (0,0,0))))
        x = torch.stack(ims).to(device)
        if device == "cuda":
            with torch.amp.autocast('cuda', dtype=torch.float16):
                feats.append(clipL.encode_image(x))
        else:
            feats.append(clipL.encode_image(x))
    v = torch.cat(feats, dim=0)
    return (v / v.norm(dim=-1, keepdim=True)).float()

text_feats = enc_text_L(all_summaries)
flat_paths = [p for paths in first_stage for p in paths]
uniq_paths = list(dict.fromkeys(flat_paths))
path_to_idx = {p:i for i,p in enumerate(uniq_paths)}
uniq_feats = enc_imgs_L(uniq_paths)

all_images = []
for i, paths in enumerate(first_stage):
    paths = [p for p in paths if isinstance(p, str) and os.path.exists(p)]
    if not paths:
        all_images.append([]); continue
    idxs = [path_to_idx[p] for p in paths]
    v = uniq_feats[idxs]
    t = text_feats[i:i+1]
    sims = (t @ v.T).squeeze(0).cpu().numpy()
    order = sims.argsort()[::-1]
    picked = [paths[j] for j in order if sims[j] >= ALIGN_FLOOR][:FINAL_K]
    if len(picked) < FINAL_K:
        used = set(picked)
        for j in order:
            if paths[j] not in used:
                picked.append(paths[j]); used.add(paths[j])
            if len(picked) == FINAL_K: break
    if not picked and paths:
        picked = [paths[order[0]]]
    all_images.append(picked)

print(f"Prepared {len(all_summaries)} summaries and images for {DATA_CSV_PATH}. N

```

➡ Rows to process: 5
 Summarizing (new CSV): 100%|██████████| 1/1 [00:03<00:00, 3.69s/it]
 Summaries ready: 5
 Prepared 5 summaries and images for new_test_data.csv. Now re-run evaluation (

--- Robust ref/candidate alignment for new_test_data.csv ---

```
import os, numpy as np, pandas as pd
from rouge_score import rouge_scorer
import bert_score
from PIL import Image
import imagehash
```

1) Load CSV from either location

```
csv_path = "new_test_data.csv"
if not os.path.exists(csv_path) and os.path.exists("/mnt/data/new_test_data.csv"):
    csv_path = "/mnt/data/new_test_data.csv"
```

```
df_new = pd.read_csv(csv_path)
```

2) Column detection + same filtering used for summarization

```
name_map = {c.lower(): c for c in df_new.columns}
TEXT_COL = (name_map.get("text") or name_map.get("content") or name_map.get("article") or
            name_map.get("body") or name_map.get("document") or name_map.get("story"))
assert TEXT_COL, f"No obvious text column found. Columns: {list(df_new.columns)}"
```

```
rows_eval = df_new[df_new[TEXT_COL].astype(str).str.strip().ne("")] .reset_index(drop=True)
```

3) Build refs from this exact slice

```
def first_k_sentences(text, k=4):
    s = str(text).replace("\n", " ")
    parts = [p.strip() for p in s.split('.') if p.strip()]
    return '. '.join(parts[:k]) + ('.' if parts[:k] else '')
```

Prefer gold summaries if present & non-empty; else use first-k sentences

```
if ('summary' in name_map) and rows_eval[name_map['summary']].notna().any():
    refs = rows_eval[name_map['summary']].astype(str).tolist()
else:
    refs = [first_k_sentences(t) for t in rows_eval[TEXT_COL].astype(str).tolist()]
```

4) Align lengths with your generated summaries

```
assert 'all_summaries' in globals(), "Run the summarization step first to define all_summaries"
cands = list(map(str, all_summaries))
```

```

m = min(len(cands), len(refs))
if len(cands) != len(refs):
    print(f"[warn] len(cands)={len(cands)} != len(refs)={len(refs)} → truncating |
cands, refs = cands[:m], refs[:m]

# 5) Quick sanity
print(f"Scoring {len(cands)} items | text col = {TEXT_COL!r}")

# 6) Lengths
src_len_words = [len(s.split()) for s in rows_eval[TEXT_COL].astype(str).tolist()]
sum_len_words = [len(s.split()) for s in cands]
print("\nLength (words)")
print(f"articles    : {np.mean(src_len_words):.1f} ± {np.std(src_len_words):.1f} (|
print(f"summaries   : {np.mean(sum_len_words):.1f} ± {np.std(sum_len_words):.1f} (|

# 7) ROUGE-L, BERTScore
scorer = rouge_scorer.RougeScorer(['rougeL'], use_stemmer=True)
rougeL = [scorer.score(r, h)['rougeL'].fmeasure for r, h in zip(refs, cands)]

P, R, F1 = bert_score.score(cands=cands, refs=refs, lang='en', rescale_with_base)
berts = F1.tolist()

def repetition_ratio(s: str, ngram=3):
    toks = s.lower().split()
    if len(toks) < ngram: return 0.0
    ngrams = [' '.join(toks[i:i+ngram]) for i in range(len(toks)-ngram+1)]
    return 1 - (len(set(ngrams)) / len(ngrams))

rep3 = [repetition_ratio(s, 3) for s in cands]

summary_df = pd.DataFrame({
    "rougeL": rougeL,
    "bertscore": berts,
    "len_words": sum_len_words,
    "rep_trigram": rep3
})
print("\nSummary metrics (mean ± sd)")
for col in ["rougeL", "bertscore", "len_words", "rep_trigram"]:
    print(f"{col:12s}: {summary_df[col].mean():.3f} ± {summary_df[col].std():.3f}")

if 'all_images' in globals():
    normalized_images = [str(p) for p in lst if isinstance(p, str) and os.path.e
    all_img_paths = [p for paths in normalized_images for p in paths]
    if not all_img_paths:
        print("\nNo images to evaluate.")

```

```
else:
```

```
    pass
```

```
➡ Scoring 5 items | text col = 'text'
```

```
Length (words)
```

```
articles    : 583.4 ± 172.6 (median 666.0)
```

```
summaries   : 54.2 ± 3.9 (median 54.0)
```

```
Some weights of RobertaModel were not initialized from the model checkpoint at
You should probably TRAIN this model on a down-stream task to be able to use :
```

```
Summary metrics (mean ± sd)
```

```
rougeL      : 0.193 ± 0.117
```

```
bertscore   : 0.099 ± 0.122
```

```
len_words   : 54.200 ± 4.324
```

```
rep_trigram : 0.000 ± 0.000
```

```
# Polished PDF for new_test_data.csv (uses all_summaries/all_images just computed)
!pip -q install reportlab
```

```
from reportlab.lib.pagesizes import A4
from reportlab.pdfgen import canvas
from reportlab.lib.units import cm
from reportlab.lib.styles import getSampleStyleSheet
from reportlab.platypus import Paragraph
from reportlab.lib.enums import TA_LEFT
from reportlab.lib import utils
from textwrap import shorten
from PIL import Image, UnidentifiedImageError
import os, shutil
```

```
# ---- Page geometry ----
```

```
PAGE_W, PAGE_H = A4
```

```
MARGIN = 1.2 * cm
```

```
CONTENT_W = PAGE_W - 2 * MARGIN
```

```
def _wrap_paragraph(text, font_size=11, leading=14):
    styles = getSampleStyleSheet()
    st = styles["Normal"]
    st.fontName = "Helvetica"
    st.fontSize = font_size
    st.leading = leading
    st.alignment = TA_LEFT
```



```

    return Paragraph(text.replace("\n", " ").strip(), st)

def _draw_paragraph(c, p: Paragraph, x, y, max_w):
    w, h = p.wrap(max_w, PAGE_H)
    p.drawOn(c, x, y - h)
    return h

def _pil_to_tmp(img: Image.Image, max_w_px=1600):
    ratio = min(1.0, max_w_px / max(img.size))
    if ratio < 1.0:
        img = img.resize((int(img.width * ratio), int(img.height * ratio)), Image
        path = f"/content/_tmp_{id(img)}.jpg"
        img.save(path, "JPEG", quality=88, optimize=True)
    return path

def _draw_image_row_paths(c, paths, x, y, max_w, height=7*cm, gap=0.5*cm):
    # paths: list[str]
    paths = [p for p in paths if isinstance(p, str) and os.path.exists(p)][:3]
    if not paths: return 0
    cols = len(paths)
    cell_w = (max_w - (cols - 1) * gap) / cols
    max_h = 0
    for ci, p in enumerate(paths):
        try:
            iw, ih = utils.ImageReader(p).getSize()
        except Exception:
            try:
                im = Image.open(p).convert("RGB")
                p = _pil_to_tmp(im)
                iw, ih = utils.ImageReader(p).getSize()
            except (UnidentifiedImageError, FileNotFoundError, OSError):
                continue
        scale = min(cell_w / iw, height / ih)
        w, h = iw * scale, ih * scale
        xi = x + ci * (cell_w + gap) + (cell_w - w) / 2
        yi = y - h
        c.drawImage(p, xi, yi, width=w, height=h)
        max_h = max(max_h, h)
    return max_h

def _draw_image_row_pils(c, imgs, x, y, max_w, height=7*cm, gap=0.5*cm):
    # imgs: list[PIL.Image.Image]
    imgs = imgs[:3]
    if not imgs: return 0
    paths = []

```

```

    for im in imgs:
        try:
            paths.append(_pil_to_tmp(im))
        except Exception:
            pass
    return _draw_image_row_paths(c, paths, x, y, max_w, height=height, gap=gap)

def _footer(c, idx):
    c.setFont("Helvetica", 9)
    c.drawRightString(PAGE_W - MARGIN, 0.8 * cm, f"Page {idx}")

def build_pdf(sources, summaries, images_per_item, out_path):
    c = canvas.Canvas(out_path, pagesize=A4)
    for idx, (src, summ, imgs) in enumerate(zip(sources, summaries, images_per_item)):
        c.setTitle("Automated News Summary & Images")
        c.setAuthor("Team42 APS360")

        # Header
        c.setFont("Helvetica-Bold", 15)
        c.drawString(MARGIN, PAGE_H - MARGIN, f"News Summary {idx}")

        # Source preview
        preview = shorten(str(src), width=600, placeholder="...")
        p1 = _wrap_paragraph(f"<b>Source (preview):</b> {preview}", font_size=10,
                             top_y = PAGE_H - MARGIN - 1.0 * cm)
        h1 = _draw_paragraph(c, p1, MARGIN, top_y, CONTENT_W)

        # Summary block
        p2 = _wrap_paragraph(f"<b>Model Summary:</b> {summ}", font_size=12, leading=1.2)
        h2 = _draw_paragraph(c, p2, MARGIN, top_y - h1 - 0.6 * cm, CONTENT_W)

        # Images row
        y_start = top_y - h1 - h2 - 1.0 * cm
        if isinstance(imgs, list) and imgs:
            if isinstance(imgs[0], Image.Image):
                _draw_image_row_pils(c, imgs, MARGIN, y_start, CONTENT_W, height=height)
            elif isinstance(imgs[0], str):
                _draw_image_row_paths(c, imgs, MARGIN, y_start, CONTENT_W, height=height)

        _footer(c, idx)
        c.showPage()
    c.save()

```


```
sources_for_pdf = rows[TEXT_COL].astype(str).tolist()[len(all_summaries):]
```

```
imgs_for_pdf = all_images[:len(sources_for_pdf)]

pdf_path = "/content/new_test_data_polished.pdf"
build_pdf(sources_for_pdf, all_summaries, imgs_for_pdf, pdf_path)

# Optional: save to Drive
OUTPUTS_DIR = "/content/drive/MyDrive/APS360_Pipeline_Outputs"
os.makedirs(OUTPUTS_DIR, exist_ok=True)
final_path = f"{OUTPUTS_DIR}/new_test_data_polished.pdf"
shutil.copy(pdf_path, final_path)

print("Saved:", final_path)
```

 Saved: /content/drive/MyDrive/APS360_Pipeline_Outputs/new_test_data_polished.pdf